

# What happens for inequality constraints?

- ★ minimize  $f(x)$  such that  $g_1(x) \leq 0, \dots, g_k(x) \leq 0$ .
- ★ not all inequality constraints play the same role: at the point  $x$  the constraint  $i$  is said to be **active** if  $g_i(x) = 0$ .
- ★ if a constraint  $g_i$  (where  $g_i$  is  $C^1$ ) is inactive at a minimizer  $x^*$  then  $g_i(x) < 0$  in a neighborhood of  $x^*$
- ★ if  $x^*$  is a minimizer of  $f(x)$  under the constraints  $g_i$  and  $g_i(x^*) < 0$  then  $g_i$  does not impose any restriction on  $f$  locally: **ignoring it produces the same result (locally)**
- ★ equality constraints generally produced **surfaces** while inequality constraints can just give **bunded regions of  $\mathbb{R}^n$** .

# Qualification of constraints

- ★ denote by  $I(x) = \{i \in \{1, \dots, k\} : g_i(x) = 0\}$  be the indices of **active constraints** at  $x$
- ★ we say that the constraints are **qualified at  $x$**  if the gradients  $(\nabla g_i(x))_{i \in I(x)}$  are linearly independent!
- ★ geometrically, as in the equality constraints case, if the constraints are qualified at  $x$  then we may define a proper **tangent space** using the family  $(\nabla g_i(x))_{i \in I(x)}$
- ★ **Special case:** if all  $g_i$  are **affine constraints** then they are automatically qualified. Why?
  - in this case the constraints also define the tangent space themselves
  - the linear independence of the gradients at a point  $x$  is equivalent to **the removal of redundant constraints**

## Theorem 4

Let  $x^*$  be a local minimizer for the inequality constrained problem

$$\min_{g(x) \leq 0} f(x)$$

and suppose that the constraints are qualified at  $x^*$ . Then the following affirmations are true:

- There exists a uniquely defined vector of **Lagrange multipliers**  $\lambda_i^* \geq 0$ ,  $i = 1, \dots, k$  such that

$$\nabla f(x^*) + \sum_{i=1}^m \lambda_i^* \nabla g_i(x^*) = 0.$$

- Moreover, if  $g_i(x^*) < 0$  then  $\lambda_i = 0$ , also called the **complementary slackness** relations. Equivalent formulation:  $\lambda_i g_i(x^*) = 0$ .

★ why are Lagrange multipliers non-negative in this case?  $x^*$  would like to "get out of the constraints" to increase the value of  $f$

★ if  $x^*$  is an interior point for  $g(x) \leq 0$  then simply  $\nabla f(x^*) = 0$

# Example: qualification of constraints

Consider the set

$$K = \{x = (x_1, x_2) \in \mathbb{R}^2 : -x_1 \leq 0, -x_2 \leq 0, -(1 - x_1)^3 + x_2 \leq 0\}.$$

★ Maximize  $J(x) = x_1 + x_2$  for  $x \in K$ .

★ making a drawing we find that immediately that the solutions are  $(0, 1)$  and  $(1, 0)$ .

★ let's check if we can write the optimality condition at the two points:

- $(1, 0)$ : constraints not qualified: unable to write the opt. cond
- $(0, 1)$ : constraints qualified: the optimality condition can be written!

# The Lagrangian - inequality case

- ★ the optimality conditions obtained involve the gradient of the objective function and the constraints.
- ★ the optimality condition can be written as the gradient of a function combining the objective and the constraints called the **Lagrangian**:  $\mathcal{L} : \mathbb{R}^n \times \mathbb{R}_+^m$

$$\mathcal{L}(x, \lambda) = f(x) + \sum_{i=1}^k \lambda_i g_i(x) = f(x) + \lambda \cdot g(x).$$

- ★ if  $x^*$  is a local minimum of  $f$  on the set  $\{g(x) \leq 0\}$  then the optimality condition tells us that there exists  $\lambda^* \in \mathbb{R}_+^m$  such that

$$\frac{\partial \mathcal{L}}{\partial x}(x^*, \lambda^*) = 0 \text{ and } \frac{\partial \mathcal{L}}{\partial \lambda}(x^*, \lambda^*) = 0$$

- ★ moreover,  $\sup_{\lambda \in \mathbb{R}_+^m} \mathcal{L}(x, \lambda) = \begin{cases} f(x) & \text{if } g(x) \leq 0 \\ +\infty & \text{otherwise} \end{cases}$  which gives

$$\min_{g(x) \leq 0} f(x) = \min_{x \in \mathbb{R}^n} \sup_{\lambda \in \mathbb{R}_+^m} \mathcal{L}(x, \lambda).$$

- ★ the minimizer of  $f$  becomes a **saddle point for the Lagrangian**

# Come back to the optimal can problem

- ★ Area of the can (to be minimized):  $A(h, r) = 2\pi r^2 + 2\pi rh$
- ★ Volume of the can (constraint):  $V(h, r) = \pi r^2 h$
- ★ finally we obtain the problem

$$\min_{V(h,r) \geq c} A(h, r).$$

- ★ **the constraint will be active!**
- ★ write the optimality condition: find  $r$  and  $h$  in terms of  $\lambda$  and finish!
- ★ in the end we find that the optimal can will have the height  $h$  equal to two times its radius  $r$ .
- ★ find now the optimal cup: only one of the two ends is filled with material!

# Saddle points

## Definition 5

We say that  $(u, p) \in U \times P$  is a saddle point of  $\mathcal{L}$  on  $U \times P$  if

$$\forall q \in P \quad \mathcal{L}(u, q) \leq \mathcal{L}(u, p) \leq \mathcal{L}(v, p) \quad \forall v \in U$$

★ when fixing  $p$ :  $v \mapsto \mathcal{L}(v, p)$  is minimal for  $v = u$

★ when fixing  $u$ :  $q \mapsto \mathcal{L}(u, q)$  is minimal for  $q = p$

★ If  $J$  is the objective function and  $F$  defines the constraint set  $K$  (equality or inequality) then a saddle point  $(u, p)$  for the Lagrangian

$$\mathcal{L}(v, q) = J(v) + q \cdot F(v)$$

verifies that  $u$  is a minimum of  $J$  on  $K$ .

★ moreover, if the Lagrangian is defined on an open neighborhood  $U$  of the constraint set  $K$  then we also recover the optimality condition

$$\nabla J(u) + \sum_{i=1}^m p_i \nabla F_i(u) = 0.$$

# Sufficient conditions

- ★ two options: go to the second order or use convexity
- ★ it is not enough to look at the second order approximation of  $f$  on the tangent space! The **curvature of the constraint** can also play a role.
- ★ the correct way is to look at the Hessian of the Lagrangian with respect to  $x$ , **reduced to the tangent space**!
- ★ in the convex case, for **inequality constraints** things are a little bit easier!
- ★ why only for inequality constraints? Imagine that equality constraints can produce **curved surfaces** and the only way to have convexity there is if they are flat!
- ★ why the choice  $g_i(x) \leq 0$  as the definition of inequality constraints? Because if all  $g_i$  are convex functions then

$$K = \{x : g_i(x) \leq 0, i = 1, \dots, k\} \text{ is a convex set.}$$



## Theorem 6 (Kuhn-Tucker)

Suppose that the functions  $f, g_i$ ,  $i = 1, \dots, k$  are  $C^1$  and convex. Define  $K$  as the set  $K = \{x : g_i(x) \leq 0\}$  and introduce the Lagrangian

$$\mathcal{L}(v, q) = f(v) + q \cdot g(v), \quad v \in \mathbb{R}^n, q \in \mathbb{R}_+^k.$$

Let  $x^*$  be a point of  $K$  where the constraints are qualified. Then the following are equivalent:

- $x^*$  is a global minimum of  $f$  on  $K$
- there exists  $\lambda^* \in \mathbb{R}^m$  such that  $(x^*, \lambda^*)$  is a saddle point for the Lagrangian
- $g(x^*) \leq 0$ ,  $\lambda^* \geq 0$ ,  $\lambda^* \cdot F(x^*) = 0$ ,  $\nabla f(x^*) + \sum_{i=1}^k \lambda_i^* \nabla g_i(x^*) = 0$ .

★ why the reverse implication works? When  $q \geq 0$  the Lagrangian

$$\mathcal{L}(v, q) = f(v) + q \cdot g(v), \quad v \in \mathbb{R}^n, q \in \mathbb{R}_+^k$$

is convex when  $f$  and  $g = (g_i)$  are convex!

★ particular case: **affine equalities!** convex and qualified!

# Handle the constraints numerically

★ we already saw two methods:

- projected gradient algorithm:

$$x_{i+1} = \text{Proj}_K(x_i - t\nabla f(x_i))$$

- penalization: include the constraint  $\{g = 0\}$  in the objective

$$\min f(x) + \frac{1}{\varepsilon} g(x)^2$$

★ we saw that the projection is not explicit in most cases! In the meantime we learned how to solve non-linear equations. Imagine the following algorithm:

- Compute  $x_i$  and the projection  $d_i$  of  $-\nabla f(x_i)$  on the tangent space (orthogonal of  $(\nabla g_j(x_i))$ )
- advance in the direction of  $d_i$ :  $x_{i+1} = x_i + \gamma_i d_i$
- project  $x_{i+1}$  on the set of constraints using the Newton method

# Conclusion on Lagrange multipliers

- we may obtain necessary optimality conditions involving equality and inequality constraints: the gradient of  $f$  is a linear combination of the gradients of the constraints
- the gradients of the constraints need to be linearly independent at the optimum: proper definition of the tangent space!
- for inequality constraints only the active constraints come into play in the optimality condition
- sufficient conditions can be found in the convex case: Kuhn-Tucker theorem
- the theory gives new ways to handle constraints numerically

# Constrained optimization

- General theoretical and practical aspects
- A quick intro to linear programming

★ maximizing or minimizing a linear function subject to linear constraints!

★ Example:

$$\max(x_1 + x_2)$$


such that  $x_1 \geq 0$ ,  $x_2 \geq 0$  and

$$\begin{array}{rcl} x_1 + 2x_2 & \leq & 5 \\ 5x_1 + 2x_2 & \leq & 11 \\ -2x_1 + x_2 & \leq & 1 \end{array}$$

★ we have some **non-negativity constraints** and the **main constraints**

# Geometric solution

★ in dimension 2 we can solve the problem by plotting the objective function on the admissible set determined by the constraints!



★ observe that in this case the solution is situated at the intersection of the lines  $5x_1 + 2x_2 = 11$  and  $x_1 + 2x_2 = 5$ .

# Theoretical observations

- ★ the gradient of  $f(x_1, x_2) = x_1 + x_2$  is  $(1, 1)$ : it is constant and never zero!
- ★ the set  $K$  determined by the linear constraints is convex
- ★ the minimum or maximum cannot be attained in the interior of  $K$ , since  $\nabla f(x) \neq 0$ !
- ★ the optimal value is on the boundary of  $K$ . Moreover there exists a vertex of the polygon where it can be found! Why?
  - start at a point  $x_0$  inside  $K$  go against the gradient till you meet an edge
  - if the function is constant along an edge then the gradient of the function and the constraint are collinear at that point: Kuhn-Tucker Theorem says that we reached the solution!
  - otherwise, follow the direction where the function decreases till reaching a vertex. Then go to the next edge and repeat the previous reasoning.
  - the process will finish: finite number of edges!
- ★ same reasoning can be applied in higher dimensions: follow the anti-gradient direction till it is collinear to the gradient of the constraint or no further decrease is possible along further facets!

★ **The Standard Maximum Problem:** Maximize  $\mathbf{c}^t \mathbf{x} = c_1 x_1 + \dots + c_n x_n$   
subject to the constraints

$$\begin{array}{rcl} a_{11}x_1 + \dots + a_{1n}x_n & \leq & b_1 \\ \vdots & & \\ a_{m1}x_1 + \dots + a_{mn}x_n & \leq & b_m \end{array} \quad \text{or } \mathbf{Ax} \leq \mathbf{b}$$

and  $x_1 \geq 0, x_2 \geq 0, \dots, x_n \geq 0$  or  $\mathbf{x} \geq 0$

★ **The Standard Minimum Problem:** Minimize  $\mathbf{y}^t \mathbf{b} = y_1 b_1 + \dots + y_m b_m$   
subject to the constraints

$$\begin{array}{rcl} a_{11}y_1 + \dots + a_{1m}y_m & \geq & c_1 \\ \vdots & & \\ a_{n1}y_1 + \dots + a_{nm}y_m & \geq & c_n \end{array} \quad \text{or } \mathbf{y}^T \mathbf{A} \geq \mathbf{c}^T$$

and  $y_1 \geq 0, y_2 \geq 0, \dots, y_m \geq 0$  or  $\mathbf{y} \geq 0$



# Example 1

## The Transportation Problem

- ★ There are  $I$  production sites  $P_1, \dots, P_I$  which supply a product and  $J$  markets  $M_1, \dots, M_J$  to which the product is shipped.
- ★ the site  $P_i$  contains  $s_i$  products and the market  $M_j$  must receive  $r_j$  products.
- ★ the cost of transportation from  $P_i$  to  $M_j$  is  $b_{ij}$
- ★ the objective is to **minimize the transportation cost while meeting the market requirements!**
- ★ denote by  $y_{ij}$  the quantity transported from  $P_i$  to  $M_j$ . Then the cost is

$$\sum_{i=1}^I \sum_{j=1}^J y_{ij} b_{ij}$$

- ★ the constraints are

$$\sum_{j=1}^J y_{ij} \leq s_i \text{ and } \sum_{i=1}^I y_{ij} \geq r_j.$$

## Example 2

### The Optimal Assignment Problem

★ There are  $I$  persons available for  $J$  jobs. The "value" of person  $i$  working 1 day at job  $j$  is  $a_{ij}$ .

★ **Objective:** Maximize the total "value"

★ the variables are  $x_{ij}$ : the proportion of person  $i$ 's time spent on job  $j$

★ the constraints are  $x_{ij} \geq 0$

$$\sum_{j=1}^J x_{ij} \leq 1, i = 1, \dots, I \text{ and } \sum_{i=1}^I x_{ij} \leq 1, j = 1, \dots, J \leq 1$$

- can't spend a negative amount of time at a job
- a person can't spend more than 100% of its time
- no more than one person working on a job

# Some Terminology

- ★ a point is said to be **feasible** if it verifies all the constraints
- ★ the set of **feasible points** is the **constraint set**
- ★ a linear programming problem is **feasible** if the constraint set is non-empty. If this is not the case then the problem is **infeasible**
- ★ every problem involving the minimization of a linear function under linear constraints can be put into standard form
  - you can change a " $\geq$ " inequality into " $\leq$ " by **changing the signs of the coefficients**
  - if a variable  $x_i$  has no sign restriction, **write it as the difference of two new positive variables**  $x_i = u_i - v_i$ ,  $u_i, v_i \geq 0$
- ★ it is possible to pass from **inequality constraints** to **equality constraints** (and the other way around)
  - $Ax = b$  is equivalent to  $Ax \leq b$  and  $Ax \geq b$
  - If  $Ax \leq b$  then add some **slack variables**  $u \geq 0$  such that  $Ax + u = b$

## Definition 7

The dual of the standard maximum problem

$$\begin{cases} \max \mathbf{c}^T \mathbf{x} \\ \text{s.t. } A\mathbf{x} \leq \mathbf{b} \text{ and } \mathbf{x} \geq 0 \end{cases}$$

is the standard minimum problem

$$\begin{cases} \min \mathbf{y}^T \mathbf{b} \\ \text{s.t. } \mathbf{y}^T A \geq \mathbf{c}^T \text{ and } \mathbf{y} \geq 0 \end{cases} \quad .$$

# Example

★ consider the problem

$$\begin{array}{ll}\text{maximize} & x_1 + x_2 \\ \text{such that} & \mathbf{x} \geq 0 \\ & x_1 + 2x_2 \leq 5 \\ & 5x_1 + 2x_2 \leq 11 \\ & -2x_1 + x_2 \leq 1\end{array}$$

★ the dual problem is

$$\begin{array}{ll}\text{minimize} & 5y_1 + 11y_2 + y_3 \\ \text{such that} & \mathbf{y} \geq 0 \\ & y_1 + 5y_2 - 2y_3 \geq 1 \\ & 2y_1 + 2y_2 + y_3 \geq 1\end{array}$$

# Relation between dual problems

## Theorem 8

*If  $\mathbf{x}$  is feasible for the standard maximum problem and  $\mathbf{y}$  is feasible for the dual problem then*

$$\mathbf{c}^T \mathbf{x} \leq \mathbf{y}^T \mathbf{b}.$$

★ The proof is straightforward:

$$\mathbf{c}^T \mathbf{x} \leq \mathbf{y}^T A \mathbf{x} \leq \mathbf{y}^T \mathbf{b}.$$

★ important consequences:

- if the standard maximum problem and its dual are both feasible, they are **bounded feasible**: the optimal values are finite!
- If there exist feasible  $\mathbf{x}^*$  and  $\mathbf{y}^*$  for the standard maximum problem and its dual such that  $\mathbf{c}^T \mathbf{x}^* = \mathbf{y}^{*T} \mathbf{b}$  then **both are optimal for their respective problems**!

## Theorem 9 (Duality)

*If a standard linear programming problem is **bounded feasible** then so is its dual, their optimal values are equal and there exist optimal solutions for both problems.*

# Solve LP problems numerically

- ★ the **simplex algorithm**: travel along vertices of the set defined by the constraints until no decrease is possible
- ★ work with the matrix  $A$  and with vectors  $\mathbf{b}$  and  $\mathbf{c}$  and modify them using **pivot rules**: similar to the ones used when solving linear systems
- ★ exploit the connection between the standard formulation and its dual
- ★ things get more complicated when we restrict the variables to be **integers**. This gives rise to **integer programming**!
- ★ algorithms solving the main types of LP problems are implemented in various Python packages: `scipy.optimize.linprog`, `pulp`.

# The simplex algorithm

- ★ bring the problem to the case of equality constraints using **slack variables**

$$\sum_{j=1}^n a_{ij}x_j \leq b_i \iff \sum_{j=1}^n a_{ij}x_j + s_i = b_i, s_i \geq 0$$

- ★ any **free variable**  $x_j \in \mathbb{R}$  should be replaced with  $u_j - v_j$  with  $u_j, v_j \geq 0$
- ★ now we can solve

$$\begin{array}{ll}\text{maximize} & \mathbf{c}^T \mathbf{x} \\ \text{subject to} & A\mathbf{x} = \mathbf{b} \\ & \mathbf{x} \geq 0\end{array}$$

- ★ start from the origin  $\mathbf{x} = 0$  and **go through the vertices of the polytype**  $A\mathbf{x} = \mathbf{b}$
- ★ at each step perform an operation similar to the **Gauss elimination**
- ★ Possible issues: **cycling**, **numerical instabilities**.



# Practical Example 1

★ the first example of a **standard maximum problem**

$$\max(x_1 + x_2)$$

such that  $x_1 \geq 0$ ,  $x_2 \geq 0$  and

$$\begin{aligned}x_1 + 2x_2 &\leq 5 \\5x_1 + 2x_2 &\leq 11 \\-2x_1 + x_2 &\leq 1\end{aligned}$$

★ we saw geometrically that the solution should be the intersection of  $x_1 + 2x_2 = 5$  and  $5x_1 + 2x_2 = 11$

```
scipy.optimize.linprog(c, A_ub=None, b_ub=None, A_eq=None,
                        b_eq=None, bounds=None, method='simplex',
                        callback=None, options=None)
```

## Practical Example 2

★ An optimal assignment problem:  $n$

	Job 1	Job 2	Job 3
Person 1	100€	120€	80€
Person 2	150€	110€	120€
Person 3	90€	80€	110€

★ assign Person  $i$  to Job  $j$  in order to minimize the total cost!

★ we can model the situation as an LP problem with 9 variables:  $x_{ij} = 1$  if and only if Person  $i$  has job  $j$ ,  $1 \leq i, j \leq 3$

★ the constraints are as follows:

- $\sum_{i=1}^3 x_{ij} = 1$ : exactly one Person for Job  $j$
- $\sum_{j=1}^3 x_{ij} = 1$ : exactly one Job for Person  $i$

★ we should also impose that  $x_i \in \{0, 1\}$ : no fractional jobs, but we'll neglect this condition and just suppose  $x_i \geq 0$ .

★ the cost is just

$$\sum_{1 \leq i, j \leq 3} c_{ij} x_{ij}$$

# Find the LP parameters

★ let's look at the matrix of the problem: 9 variables and 6 constraints!

$$A = \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$

★ the matrix  $c_{ij}$  is given by the table shown previously: the cost of every person per function ★ the vector  $b$  is equal to 1 on every component

★ the solution is made of zeros and ones, without imposing this...

★ this phenomenon always happens: if  $A$  is a **totally unimodular matrix** and  $b$  is made of integers then  $Ax \leq b$  has all its vertices at points with **integer coordinates**

A matrix is **totally unimodular** if every square submatrix has determinant in the set  $\{0, 1, -1\}$ .

# Practical Example 3

★ solving a Sudoku with LP

					3		8	5
		1		2				
			5		7			
		4				1		
	9							
5							7	3
		2		1				
				4				9

- ★ Remember the rules:  $\{1, 2, 3, 4, 5, 6, 7, 8, 9\}$  should be found on every line, column and  $3 \times 3$  square
- ★ in order to make this solvable via LP a different formulation should be used!
- ★ classical idea: use binary variables

# Sudoku in Binary variables

- ★ how to represent a Sudoku puzzle using 0s and 1s?
- ★ build a 3D array  $X = (x_{ijk})$  of size  $9 \times 9 \times 9$  such that
$$x_{ijk} = 1 \text{ if and only if on position } (i, j) \text{ we have the digit } k; \text{ else } x_{ijk} = 0$$
- ★ what are the constraints in this new formulation?
  - $x_{ijk} \in \{0, 1\}$ : again to be relaxed to  $x_{ijk} \geq 0$  - 729 constraints
  - fixing  $i, j$ :  $\sum_{k=1}^9 x_{ijk} = 1$  - one number per cell - 81 constraints
  - fixing  $i, k$ :  $\sum_{j=1}^9 x_{ijk} = 1$  -  $k$  appears exactly once on line  $i$  - 81 constraints
  - fixing  $j, k$ :  $\sum_{i=1}^9 x_{ijk} = 1$  -  $k$  appears exactly once on column  $j$  - 81 constraints
  - small  $3 \times 3$  squares condition: for  $u, v \in \{0, 3, 6\}$ 
$$\sum_{i=1}^3 \sum_{j=1}^3 x_{i+u, j+v, k} = 1, \quad k = 1, \dots, 9 \text{ - 81 constraints}$$
  - the initial given information for the puzzle may be written in the form  $s_{ij} = k$  for some  $i, j, k$ . This gives the constraints  $x_{i,j,s_{ij}} = 1$ .
- ★ we are interested in finding a feasible solution: no objective function is needed!

# Solving the Sudoku

- ★ a feasible solution can be found using the simplex algorithm
- ★ sometimes we may get **non-integer results**: apparently, the constraint matrix is not always a Totally Unimodular matrix
- ★ there are LP algorithms which will return integer solutions: **integer programming**
- ★ before solving we should check that **the constraint matrix should be of maximal rank**: eliminate redundant constraints
- ★ we could also eliminate **fixed variables**: the data  $s_{ij} = k$  should eliminate all unknowns with first index  $i$ , second index  $j$  or third index  $k$ !
- ★ if the solution is unique: the algorithm **will find it**
- ★ if the solution is not unique: the algorithm will find **one of the solutions**. We may repeat with the constraint that the solution should be different than the previous one, until no other solutions are found!
- ★ check out the PuLP Python library: an example of Sudoku solver is given!

# Conclusions on LP

- minimize/maximize linear functions under linear constraints
- many practical applications from an industrial point of view!
- there exist optimizers which are **vertices of the constraint set**
- **simplex algorithm**: travel along vertices decreasing the objective function
- computational complexity: **worst case is exponential**: Klee-Minty cube
- polynomial-time **average case complexity**: most of the LP problems will be solved very fast!

# Conclusion of the course

- ★ numerical optimization (unconstrained case):
  - derivatives-free algorithms: no-regularity needed, slow convergence
  - gradient descent algorithms: linear convergence, sensitive to the condition number
  - Newton, quasi-Newton: super-linear convergence in certain cases
  - when dealing with large problems use L-BFGS
  - Conjugate Gradient: solve linear systems, better than GD
  - Gauss-Newton: useful when minimizing a **non-linear least squares function**
- ★ constrained case
  - for simple constraints: use the projected gradient algorithm
  - general smooth constraints: use the **tangential part of the gradient** and come back to the constraint set using the **Newton method**
  - other options available: SQP, etc...
  - Linear Programming: use specific techniques: the **simplex algorithm** → to be continued next year in the course dealing with Convex Optimization!



# Conclusion of the course

- know your options when looking at an optimization problem: choose the right algorithm depending on: the size of the problem, the number of variables, the regularity, the conditioning, etc.
- learn how to use existing solutions: `scipy.optimize` is a good starting point
- know how to code your own optimization algorithm if necessary: use gradients when possible, limit the number of function evaluations, choose a good stopping criterion, limit the number of iterations, etc.